

# Catalogue

1 Analyzing $\mathcal{M}_{means}$ .....	1
(a).....	1
(b) .....	2
(c).....	2
2 Implementing $\mathcal{M}_{means}$ .....	3
(a).....	3
(b) .....	4
Appendix .....	5

Dengke Chen and qvl920

## 1 Analyzing $\mathcal{M}_{means}$

(a)

Answer:

To show that  $\mathcal{M}_{means}$  satisfies  $\rho$ -zCDP given that  $\mathcal{M}'_{means}$  does, need to establish that  $\mathcal{M}_{means}$  is a post-processing of  $\mathcal{M}'_{means}$ . According to Lemma 4 (Composition), privacy guarantees are preserved under post-processing; that is, any function applied to the output of a  $\rho$ -zCDP mechanism remains  $\rho$ -zCDP.

$\mathcal{M}'_{means}$ : Outputs noisy versions of intermediate computations:

$$\{f_\ell(X) + z_\ell, g_\ell(X) + z'_\ell\}_{\ell=1}^t,$$

Here,  $f_\ell(X)$  and  $g_\ell(X)$  are functions of the input data  $X$ , and  $z_\ell$  and  $z'_\ell$  are Gaussian noise vectors.

$\mathcal{M}_{means}$ : Outputs the final cluster centers:

$$\{c_i^{(t)}\}_{i=1}^k,$$

The  $c_i^{(t)}$  are computed using the noisy sums and counts from  $\mathcal{M}'_{means}$ .

The cluster centers  $\{c_i^{(t)}\}$  in  $\mathcal{M}_{means}$  are computed using the noisy outputs from  $\mathcal{M}'_{means}$  through deterministic calculations.

Specifically, the update rule for cluster centers in step 2.(a) of  $\mathcal{M}_{means}$  is

$$c_i^{(\ell)} = \frac{1}{\max(1, n_i^{(\ell-1)})} \left( z_{\ell,i} + \sum_{j \in S_i^{(\ell-1)}} x_j \right) = \frac{1}{\max(1, n_i^{(\ell-1)})} (f_\ell(X)_i + z_{\ell,i})$$

This is a function of  $f_\ell(X) + z_\ell$  and  $n_i^{(\ell-1)}$  (which depends on  $g_{\ell-1}(X) + z'_{\ell-1}$ ).

Since  $\mathcal{M}_{means}$  computes its output by applying deterministic functions to the outputs of  $\mathcal{M}'_{means}$ , it is a post-processing of  $\mathcal{M}'_{means}$ .

Lemma 4 (Composition) tells us that post-processing a  $\rho$ -zCDP mechanism does not increase the privacy loss.

Therefore, if  $\mathcal{M}'_{means}$  satisfies  $\rho$ -zCDP, then  $\mathcal{M}_{means}$  also satisfies  $\rho$ -zCDP.

Since  $\mathcal{M}_{means}$  is a deterministic post-processing of  $\mathcal{M}'_{means}$ , and privacy guarantees are preserved under post-processing,  $\mathcal{M}_{means}$  satisfies  $\rho$ -zCDP if  $\mathcal{M}'_{means}$  does.

(b)

Answer:

$X$  and  $X'$  differ in exactly one data point, say at index  $s$ .

Since  $S_1^{(\ell-1)}$  are fixed and independent of  $X$ , the sets remain the same for both  $X$  and  $X'$ .

For each cluster  $i \in [k]$ :

$$\begin{aligned} f_\ell(X)_i &= \sum_{j \in S_i^{(\ell-1)}} x_j \\ f_\ell(X')_i &= \sum_{j \in S_i^{(\ell-1)}} x'_j \end{aligned}$$

Since  $x_j = x'_j$  for all  $j \neq s$ , the only potential difference comes from  $j = s$  if  $s \in S_1^{(\ell-1)}$ .

The difference vector  $f_\ell(X) - f_\ell(X')$  has non-zero components only in the cluster  $i$  such that  $s \in S_1^{(\ell-1)}$ :

$$[f_\ell(X) - f_\ell(X')]_i = \begin{cases} x_s - x'_s, & \text{if } s \in S_1^{(\ell-1)} \\ 0, & \text{otherwise} \end{cases}$$

Since  $\|x_s\|_2, \|x'_s\|_2$  we have:

$$\|x_s - x'_s\|_2 \leq \|x_s\|_2 + \|x'_s\|_2 \leq 2$$

Therefore, the squared  $\ell_2$ -norm of the difference is:

$$\|f_\ell(X) - f_\ell(X')\|_2^2 = \sum_{i=1}^k \|[f_\ell(X) - f_\ell(X')]_i\|_2^2 \leq 4$$

Final, taking the square root, get:  $\|f_\ell(X) - f_\ell(X')\|_2 \leq 2$

By fixing  $S_1^{(\ell-1)}$ , the difference  $f_\ell(X) - f_\ell(X')$  is non-zero only in one cluster  $i$  where  $s \in S_1^{(\ell-1)}$ , and this difference has norm at most 2. Therefore,  $\|f_\ell(X) - f_\ell(X')\|_2 \leq 2$  for neighboring  $X \sim X'$ .

(c)

Answer:

Privacy Loss per Iteration:

In each iteration  $\ell$ ,  $\mathcal{M}'_{means}$  releases:

Noisy Sum:  $f_\ell(X) + z_\ell$  where  $z_\ell \sim \mathcal{N}(0, \sigma^2)^{dk}$

Noisy Counts:  $g_\ell(X) + z'_\ell$  where  $z'_\ell \sim \mathcal{N}(0, \sigma'^2)^k$

Sensitivity Computations:

Sensitivity of  $f_\ell(X)$ :

From Problem (b),  $\|f_\ell(X) - f_\ell(X')\|_2 \leq 2$ , So,  $\Delta_f = 2$

Sensitivity of  $g_\ell(X)$ :

Changing one data point can affect cluster assignments, changing the counts in up to two clusters by 1 (one increases by 1, another decreases by 1).

The difference vector has +1 and -1 in two components, zeros elsewhere.

$$\|g_\ell(X) - g_\ell(X')\|_2 = \sqrt{2},$$

$$\text{So, } \Delta_g = \sqrt{2}$$

Applying the Gaussian Mechanism (Proposition 3):

$$\text{Privacy Loss for } f_\ell(X) + z_\ell: \rho_f = \frac{\Delta_f^2}{2\sigma^2} = \frac{2^2}{2\sigma^2} = \frac{2}{\sigma^2}$$

$$\text{Privacy Loss for } g_\ell(X) + z'_\ell: \rho_g = \frac{\Delta_g^2}{2\sigma'^2} = \frac{\sqrt{2}^2}{2\sigma'^2} = \frac{1}{\sigma'^2}$$

$$\text{Total Privacy Loss per Iteration: } \rho_{iter} = \rho_f + \rho_g = \frac{2}{\sigma^2} + \frac{1}{\sigma'^2}$$

Aggregating Over  $t$  Iterations:

Using Lemma 4 (Composition), the total privacy loss over  $t$  iterations is additive:

$$\rho_{total} = t \cdot \rho_{iter} = t(\rho_f + \rho_g) = \left(\frac{2t}{\sigma^2} + \frac{t}{\sigma'^2}\right)$$

Therefore  $\mathcal{M}'_{means}$  satisfies  $\rho$ -zCDP with  $\rho = 2t/\sigma^2 + t/\sigma'^2$ .

## 2 Implementing $\mathcal{M}_{means}$

(a)

Answer:

The link to the full code on Colab: [my code on Colab](#)

From the previous analysis, the algorithm satisfies  $\rho$ -zCDP with:  $\rho = 2t/\sigma^2 + t/\sigma'^2$

To allocate the privacy budget optimally between  $\sigma$  and  $\sigma'$ , balance the privacy loss contributed by each. Using the sensitivities of the functions:

Sensitivity of  $f_\ell(X)$ :  $\Delta_f = 2$

Sensitivity of  $g_\ell(X)$ :  $\Delta_g = \sqrt{2}$

Set:  $\frac{\Delta_f^2}{2\sigma^2} = \frac{\Delta_g^2}{2\sigma'^2} = \frac{\rho}{2t}$

Solving for  $\sigma^2$  and  $\sigma'^2$ :  $\sigma^2 = \frac{4t}{\rho}$ ,  $\sigma'^2 = \frac{2t}{\rho}$

The modified code section can be found in the file appendix.

Algorithm Steps:

Initialization:

Randomly assigns each data point to one of the  $k$  clusters.

Computes the initial noisy cluster counts  $n_i^{(0)}$ .

Iterative Updates (for  $\ell = 1$  to  $t$ ):

Cluster Center Update [Step 2.(a)]:

Computes the sum of points in each cluster.

Adds Gaussian noise  $z_{\ell,i}$  to the sum.

Updates the cluster centers  $c_1^{(\ell)}$  using the noisy counts  $n_i^{(\ell-1)}$ .

Cluster Assignment [Step 2.(b)]:

Assigns each point to the nearest cluster center.

Noisy Count Update [Step 2.(c)]:

Computes the new cluster sizes.

Adds Gaussian noise  $\mathcal{Z}'_{\ell,i}$  to obtain  $n_i^{(\ell)}$ .


Final Output [Step 3]:

Returns the cluster centers after  $t$  iterations.

(b)

Answer:

Result plot (The modified code section can be found in the file appendix.):



High Privacy (Low  $\rho$ ): At low values of  $\rho$ , the added noise is significant, leading to higher clustering costs.

Low Privacy (High  $\rho$ ): As  $\rho$  increases, less noise is added, resulting in better clustering (lower cost).

## Appendix

```
# Differentially Private k-means Algorithm
def M_means(points, k, t, rho):
    n, d = points.shape

    # Calculate sigma and sigma_prime based on rho and t
    sigma_squared = (4 * t) / rho
    sigma = np.sqrt(sigma_squared)

    sigma_prime_squared = (2 * t) / rho
    sigma_prime = np.sqrt(sigma_prime_squared)

    # Step 1: Random initialization of disjoint clusters
```

```

initial_assignment = np.random.choice(range(k), n)
cluster_indexes = [ np.where(initial_assignment == i)[0] for i in
range(k) ]
n_i_prev = [ len(cluster_indexes[i]) for i in range(k) ] #
 $n_i^{(0)}$ 

for l in range(1, t + 1):
    # Step 2(a): Update cluster centers with noise
    c_i = []
    for i in range(k):
        S_i_prev = cluster_indexes[i]
        # Sum over points in  $S_i^{(l-1)}$ 
        if len(S_i_prev) > 0:
            f_l_i = np.sum(points[S_i_prev], axis=0)
        else:
            f_l_i = np.zeros(d)
        # Add Gaussian noise
        Z_l_i = np.random.normal(0, sigma, d)
        denominator = max(1, n_i_prev[i])
        # Compute  $c_i^{(l)}$ 
        c_i_l = (Z_l_i + f_l_i) / denominator
        c_i.append(c_i_l)
    centers = np.array(c_i) # Centers for iteration l

    # Step 2(b): Assign points to the nearest cluster center
    distances_squared = np.sum((points - centers[:,
np.newaxis])**2, axis=-1)
    assignment = np.argmin(distances_squared, axis=0)
    cluster_indexes = [ np.where(assignment == i)[0] for i in
range(k) ]

    # Step 2(c): Update cluster sizes with noise
    n_i = []
    for i in range(k):
        cluster_size = len(cluster_indexes[i])
        z_prime_l_i = np.random.normal(0, sigma_prime)
        n_i_l = cluster_size + z_prime_l_i
        n_i.append(n_i_l)
    n_i_prev = n_i # Update for next iteration
    # Step 3: Output the final cluster centers
    return centers # Final cluster centers  $c_i^{(t)}$ 

```

```

"""# Plot cost as function of number of iterations"""

# Parameters
k = 5 # Number of clusters
t = 5 # Number of iterations
rho_values = np.logspace(-3, 0, num=20) # 20 values from 0.001 to 1
costs = []

# Run M_means for each rho and compute cost
for rho in rho_values:
    centers = M_means(points, k, t, rho)
    cost = compute_cost(points, centers)
    costs.append(cost)
    print(f"Rho: {rho:.4f}, Cost: {cost:.4f}")

# Plot the cost as a function of rho
fig, ax = plt.subplots()
ax.set_xlabel('ρ')
ax.set_ylabel('Cost')
ax.plot(rho_values, costs, marker='o')
ax.set_xscale('log')
ax.set_title('Cost vs. ρ')
plt.show()

```